

# Zkouška EN ústní

## ▼ Introduction to DP

### 1. Core Concepts

Term	Definition
<b>Dynamic Programming</b>	Method for solving multi-stage decision problems by breaking them into smaller subproblems. Relies on Bellman's <b>Principle of Optimality</b> .
<b>Stage</b>	Discrete time index $k = 0, 1, \dots, N$ . Decisions are made stage by stage.
<b>State <math>x_k</math></b>	Information summarizing the past that is relevant for future optimization.
<b>Control <math>u_k</math></b>	Decision chosen at stage $k$ . Must satisfy $u_k \in U_k(x_k)$ .
<b>Disturbance <math>w_k</math></b>	Random variable affecting transitions. Distribution depends only on current state & control ( <b>Markov property</b> ).
<b>System Dynamics</b>	$x_{k+1} = f_k(x_k, u_k, w_k)$
<b>Cost Function</b>	Additive form: terminal cost $g_N(x_N)$ + stage costs $g_k(x_k, u_k, w_k)$ . Optimization based on expected total cost.
<b>Policy <math>\pi = \{\mu_0, \dots, \mu_{N-1}\}</math></b>	Set of functions mapping states to decisions: $u_k = \mu_k(x_k)$ .
<b>Admissible Policy</b>	Policy where every $\mu_k(x_k) \in U_k(x_k)$ .
<b>Optimal Policy <math>\pi^*</math></b>	Minimizes expected total cost.
<b>Optimal Cost / Value Function <math>J^*(x_0)</math></b>	Minimum achievable expected cost from initial state $x_0$ .

### 2. Open-Loop vs. Closed-Loop

Mode	Description	Pros/Cons
<b>Open-loop</b>	Choose all controls $u_0, \dots, u_{N-1}$ at time 0. No adaptation to disturbances.	Simpler but suboptimal—cannot react to new information.

Mode	Description	Pros/Cons
<b>Closed-loop</b>	Choose $u_k$ at time $k$ based on current state $x_k$ .	Always $\geq$ performance; yields <b>value of information</b> .

### 3. Transition Models

#### Continuous / General Form

- $x_{k+1} = f_k(x_k, u_k, w_k)$

#### Discrete-State Transition Probabilities

- $p_{ij}(u, k) = P\{x_{k+1} = j \mid x_k = i, u_k = u\}$

Equivalent to describing dynamics via random disturbance  $w_k$ .

### 4. Expected Total Cost Under Policy

$$J_\pi(x_0) = \mathbb{E} \left[ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, \mu_k(x_k), w_k) \right]$$

Optimal cost:

$$J^*(x_0) = \min_\pi J_\pi(x_0)$$

### 5. Principle of Optimality (Bellman)

Given an optimal policy, its tail starting from any reachable state at time  $i$  is itself optimal for the subproblem from  $i$  to  $N$ .

This enables **backward induction** → foundation of the DP algorithm.

### 6. DP Backward Recursion (Finite Horizon)

From slide proposition:

1. **Initialization (terminal stage)**

$$J_N(x_N) = g_N(x_N)$$

2. **Backward recursion for  $k = N - 1, \dots, 0$**

$$J_k(x_k) = \min_{u_k \in U_k(x_k)} \mathbb{E}_{w_k} [g_k(x_k, u_k, w_k) + J_{k+1}(f_k(x_k, u_k, w_k))]$$

3. **Optimal control law**

$$\mu_k^*(x_k) = \arg \min_{u_k \in U_k(x_k)} \mathbb{E}_{w_k} [g_k(x_k, u_k, w_k) + J_{k+1}(f_k(x_k, u_k, w_k))]$$


---

## 7. Interpretation of $J_k(x_k)$

- Cost-to-go function from state  $x_k$  at time  $k$ .
  - Represents optimal expected cost for a remaining horizon of  $N - k$ .
  - Slide note: Computational effort grows with number of states → **curse of dimensionality**.
- 

## 8. Example – Inventory Control (Stochastic)

### Variables

Symbol	Description
$x_k$	Stock at start of period $k$ .
$u_k$	Units ordered at start of period $k$ .
$w_k$	Random demand in period $k$ .

### Dynamics

$$x_{k+1} = x_k + u_k - w_k$$

### Cost Components

- Holding / shortage cost:  $r(x_k)$
- Ordering cost:  $cu_k$
- Terminal cost:  $R(x_N)$

### Objective

$$\mathbb{E} \left[ R(x_N) + \sum_{k=0}^{N-1} (r(x_k) + cu_k) \right]$$

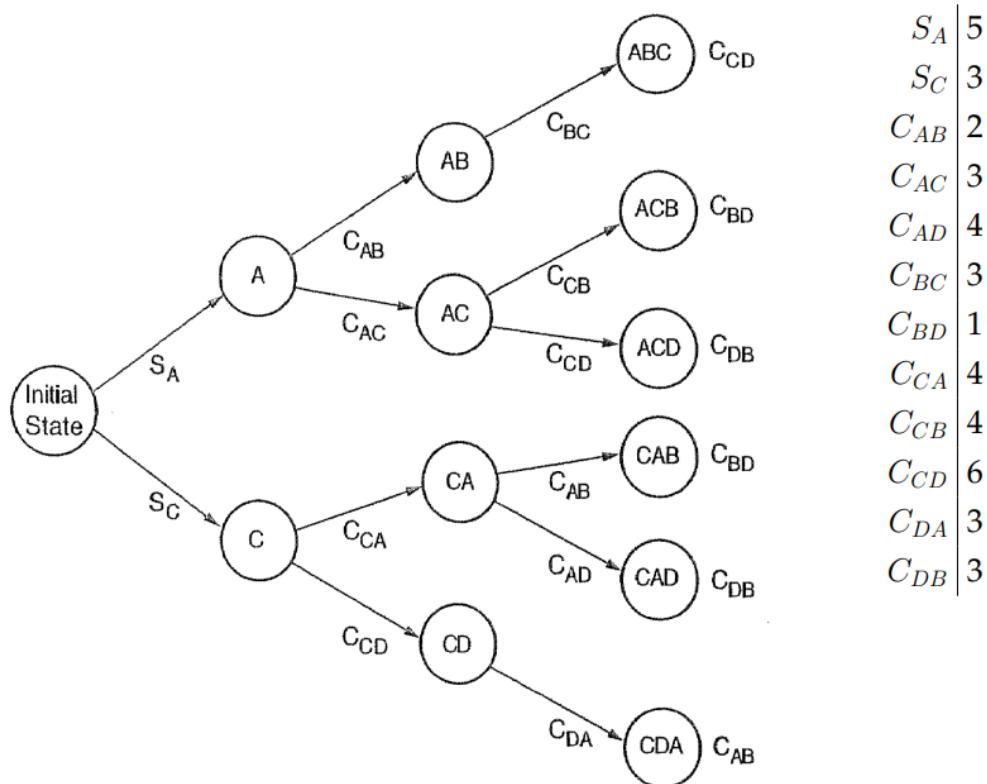
DP chooses ordering rule  $u_k = \mu_k(x_k)$ .

---

## 9. Example – Deterministic Scheduling

- Sequence operations A,B,C,D with precedence constraints.

- Each decision chooses next operation; cost is setup cost + startup cost.
- Represented as **transition graph**.



- Solved naturally via DP as a finite-state deterministic problem.

## 10. Why DP Works

- Converts a full-horizon optimization into nested smaller problems.
- Backward induction computes optimal decisions *stage by stage*.
- Optimal policy is *state-feedback* → inherently adaptive.

## Key Formulas Summary

Concept	Formula
<b>Dynamics</b>	$x_{k+1} = f_k(x_k, u_k, w_k)$
<b>Total Expected Cost</b>	$J_\pi(x_0)$ as above

Concept	Formula
<b>DP Recursion</b>	$J_k(x_k) = \min_u \mathbb{E}[g_k + J_{k+1}]$
<b>Optimal Control Law</b>	$\mu_k^*(x_k) = \arg \min_u \mathbb{E}[g_k + J_{k+1}]$
<b>Terminal Condition</b>	$J_N(x_N) = g_N(x_N)$

## ▼ Minimax

### 1. Core Idea of Minimax Control

Minimax (robust) control addresses decision-making under uncertainty  
**when no probabilistic description is available.**

Instead of assuming known probability distributions for disturbances, we only know **sets** within which disturbances may lie.

The goal is to design a policy that performs well **under the worst possible disturbance sequence.**

Formally, the objective is:

$$\min_{\pi \in \Pi} \max_{w \in W} J(\pi, w)$$

This “worst-case” approach ensures **robustness** in situations where uncertainty cannot be modeled stochastically.

### 2. Problem Structure

Disturbances are assumed to satisfy:

$$w_k \in W_k(x_k, u_k)$$

A policy is a sequence of state-feedback functions:

$$\pi = \{\mu_0, \mu_1, \dots, \mu_{N-1}\}$$

The worst-case cost associated with a policy is:

$$J_\pi(x_0) = \max_{w_k \in W_k(x_k, \mu_k(x_k))} \left[ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, \mu_k(x_k), w_k) \right]$$

This replaces the **expected cost** from stochastic DP with a **maximum cost** over all admissible disturbances.

### 3. Minimax Dynamic Programming Algorithm

Just like in classical DP, the minimax algorithm proceeds backward.

The difference is that the expectation operator is replaced by a maximization over disturbances.

## Terminal stage

$$J_N(x_N) = g_N(x_N)$$

## Backward recursion

$$J_k(x_k) = \min_{u_k \in U_k(x_k)} \max_{w_k \in W_k(x_k, u_k)} [g_k(x_k, u_k, w_k) + J_{k+1}(f_k(x_k, u_k, w_k))]$$

The controller selects the action that gives the **best performance under the worst disturbance**.

## 4. State Augmentation (Yellow Highlights)

When assumptions of the classical DP model are violated (e.g., dependence on earlier states or controls), we can reformulate the problem through **state augmentation**.

General rule:

Include in the state all information available at time  $k$  that is relevant for choosing  $u_k$ .

Examples:

### Time lags

If dynamics depend on earlier values:

$$x_{k+1} = f_k(x_k, x_{k-1}, u_k, u_{k-1}, w_k)$$

Introduce:

- $y_k = x_{k-1}$
- $s_k = u_{k-1}$

New augmented state:

$$\tilde{x}_k = (x_k, y_k, s_k)$$

**Consequence:** State space expands dramatically.

If original system had  $n$  states and  $m$  controls, augmented state may have up to:

$$n \times n \times m$$

## 5. Correlated Disturbances

If disturbances follow a correlation model:

$$w_k = \lambda w_{k-1} + \xi_k$$

Introduce an additional state variable:

$$y_k = w_{k-1}$$

New system state:

$$\tilde{x}_k = (x_k, y_k)$$

This converts correlated noise into a DP-compatible form.

## 6. Simplification When Part of the State Is Uncontrollable

If the state consists of  $(x_k, y_k)$  but only  $x_k$  can be influenced by control, computation can be reduced.

Define the reduced cost-to-go:

$$\hat{J}_k(x_k) = E_{y_k} \{ J_k(x_k, y_k) \mid x_k \}$$

This allows DP to operate over the **controllable** component only, while still yielding an optimal control law defined over the full state.

## 7. Parking Problem (Illustrative Example)

In the parking problem, each space is either **free** or **taken**, and the driver must decide whether to park or move on.

The state distinguishes:

- $(k, F)$ : space  $k$  is free
- $(k, \bar{F})$ : space  $k$  is taken
- Terminal state  $t$
- Garage with cost  $C$

The cost-to-go functions are:

$$J_k^*(F) = \begin{cases} \min[c(k), p(k+1)J_{k+1}^*(F) + (1-p(k+1))J_{k+1}^*(\bar{F})], & k < N-1 \\ \min[c(N-1), C], & k = N-1 \end{cases}$$
$$J_k^*(\bar{F}) = \begin{cases} p(k+1)J_{k+1}^*(F) + (1-p(k+1))J_{k+1}^*(\bar{F}), & k < N-1 \\ C, & k = N-1 \end{cases}$$

Define the expected cost before knowing whether a space is free:

$$\hat{J}_k = p(k) \min[c(k), \hat{J}_{k+1}] + (1-p(k))\hat{J}_{k+1}$$

**Optimal policy :**

| Park at space  $k$  if it is free and  $c(k) \leq \hat{J}_{k+1}$ .

## Key Takeaways (What to Say on an Exam)

- Minimax control is used when uncertainty **cannot be described probabilistically**.
- Disturbances are assumed to lie in **known sets**, and the objective is to minimize the **worst-case** cost.
- The dynamic programming recursion becomes a **min-max** problem instead of **min-expectation**.
- State augmentation is used to handle time lags or correlated disturbances.
- When some state components are uncontrollable, DP can be simplified by averaging over them.
- The parking problem illustrates DP with uncontrollable randomness and reduced-state formulation.

## ▼ Deterministic finite-state problems

### 1. Deterministic Systems: Key Characteristics

A deterministic system is one where each disturbance  $w_k$  can take **only one value**.

Such problems occur naturally when uncertainty is negligible or when a stochastic model is approximated by fixing disturbances to a typical value.

A fundamental property of deterministic systems is:

Feedback provides no cost advantage.

Minimizing over feedback policies  $\{\mu_0, \dots, \mu_{N-1}\}$  yields the same optimal cost as minimizing over open control sequences  $\{u_0, \dots, u_{N-1}\}$ , because future states evolve **predictably** under deterministic dynamics.

Given a policy, the state evolution is fully known:

$$x_{k+1} = f_k(x_k, \mu_k(x_k)), \quad u_k = \mu_k(x_k)$$

This distinguishes deterministic problems from stochastic ones and has computational consequences:

deterministic continuous-space problems can be solved by variational or gradient-based optimization, whereas DP remains useful when constraints or discrete structure are present.

## 2. Deterministic Finite-State Systems and the Shortest Path Interpretation

Consider a deterministic problem where each stage  $k$  has a **finite state set**  $S_k$ .

Every control  $u_k$  induces a transition:

$$x_{k+1} = f_k(x_k, u_k)$$

with associated cost  $g_k(x_k, u_k)$ .

This leads to a natural **graph representation**:

- **Nodes** represent states at each stage.
- **Arches** represent feasible transitions  $x_k \rightarrow x_{k+1}$ .
- **Arc cost** = transition cost  $g_k(x_k, u_k)$ .
- An artificial terminal node  $t$  is added, and each final state  $x_N$  has an arc  $(x_N \rightarrow t)$  with cost  $g_N(x_N)$ .

## Core Interpretation

A deterministic finite-state problem is equivalent to finding a minimum-length (shortest) path from the initial node  $s$  to the terminal node  $t$ .

A **path** is a sequence of arcs

$$(j_1, j_2), (j_2, j_3), \dots, (j_{k-1}, j_k)$$

and its **length** is the sum of arc costs.

This graphical formulation makes DP identical to a shortest-path computation.

---

## 3. DP Algorithm for Deterministic Finite-State Problems

Using the notation:

- $a_{i,j}^k$ : cost of transition at stage  $k$  from state  $i \in S_k$  to state  $j \in S_{k+1}$
- $N = gN(i) = g_N(i)_{ai,tN=gN(i)}$ : terminal cost
- $a_{i,j}=\infty$  if no control allows the transition

Dynamic programming becomes:

### Terminal stage

$$J_N(i) = a_{i,tN}, i \in S_N, N \in \mathbb{N}$$

$$J_N(i) = \infty, i \notin S_N$$

### Backward recursion

$$\begin{aligned} J_k(i) &= \min_{j \in S_{k+1}} [a_{i,j} + J_{k+1}(j)], k = 0, \dots, N-1 \\ &= \min_{j \in S_{k+1}} [a_{i,j} + J_{k+1}(j)] \quad \forall k = 0, \dots, N-1 \end{aligned}$$

$$J_k(i) = \min_{j \in S_{k+1}} [a_{i,j} + J_{k+1}(j)], k = 0, \dots, N-1$$

The optimal cost:

$$J_0(s)$$

$$J_0(s)$$

is exactly the **length of the shortest path from  $s$  to  $t$** .

## 4. Forward DP Algorithm (Special Case)

Although DP is usually backward, deterministic finite-state problems allow an equivalent **forward algorithm**.

Key observation (yellow highlight):

An optimal path from  $s$  to  $t$  is also optimal in a reverse graph, where all arcs are reversed but their lengths remain unchanged.

The forward DP computes **cost-to-arrive** from  $s$ :

$$J \sim N(j) = \min_{j \in S_1} \tilde{J}_N(j) = a_{s,j}^0, \quad j \in S_1$$

$$J \sim N(j) = \min_{j \in S_1} \tilde{J}_N(j)$$

$$J \sim k(j) = \min_{i \in S_{N-k}} [a_{i,j}^{N-k} + \tilde{J}_{k+1}(i)]$$

$$J \sim k(j) = \min_{i \in S_{N-k}} [a_{i,j}^{N-k} + \tilde{J}_{k+1}(i)]$$

The final cost satisfies:

$$J \sim 0(t) = J_0(s) = \tilde{J}_0(t) = J_0(s)$$

$$J \sim 0(t) = J_0(s)$$

Forward DP is useful in **real-time applications** where stage- $k$  data becomes available only when stage  $k$  is reached.

## 5. Shortest Path Problems as Deterministic DP

Any classical shortest-path problem can be formulated as a deterministic finite-state DP problem.

Given:

- nodes  $\{1, 2, \dots, N, t\}$ ,
- arc cost  $a_{ij}$ ,
- $a_{ij} = \infty$  if arc does not exist,

Goal: find shortest path from each node  $i$  to  $t$ .

Assumption (yellow highlight):

No cycle may have negative total length, ensuring optimal paths never need more than  $N$  moves.

Hence we can require exactly NNN moves by allowing "degenerate moves"  $i \rightarrow ii$  \to  $ii \rightarrow i$  with zero cost.

DP definition:

$J_k(i)$ =optimal cost of reaching  $t$  in  $N-k$  moves  
 $J_k(i) = \min_{j=1,\dots,N} [a_{ij} + J_{k+1}(j)]$ ,  $k=0,\dots,N-2$

$J_k(i)$ =optimal cost of reaching  $t$  in  $N-k$  moves

DP equations:

$J_k(i) = \min_{j=1,\dots,N} [a_{ij} + J_{k+1}(j)]$ ,  $k=0,\dots,N-2$

$J_{N-1}(i) = \min_{t=1,\dots,N} [a_{it}]$

$J_{N-1}(i) = a_{it}$

If degenerate moves appear, it simply means the actual shortest path uses fewer than NNN transitions.

---

## Key Takeaways (What to Say on the Exam)

- Deterministic systems have **predictable** state evolution; feedback offers **no advantage**.
- Finite-state deterministic DP problems map naturally to **shortest-path problems**.
- DP recursions compute shortest paths via backward or forward algorithms.
- Graph representation: **states = nodes, controls = arcs, transition cost = arc length**.
- DP and shortest-path formulations are fully equivalent.
- Forward DP is useful when information is revealed sequentially.
- Any shortest path problem can be converted into deterministic DP under the assumption of **no negative cycles**.

## ▼ Shortest path methods